Case Study 2
AKSTA Statistical Computing

Tatzberger Jonas, Rasser Thomas, Grübling Alexander

03.05.2024

Exercises

1. Data Import

a.

Load in R the following data sets which you can find in TUWEL. For each data set, ensure that missing values are read in properly, that column names are unambiguous. Each data set should contain at the end only two columns: country and the variable.

Answer

```
check_number_of_rows <- function(expected_rowcount, expected_colcount, given_tibble) {
  real_rowcount = nrow(given_tibble)
  real_colcount = ncol(given_tibble)

  result <- assert_that(real_rowcount == expected_rowcount,
              msg = paste0("There should be ",
                           expected_rowcount,
                           " rows instead of ",
                           real_rowcount))

  result <- assert_that(real_colcount == expected_colcount,
              msg = paste0("There should be ",
                           real_colcount,
                           " columns instead of ",
                           expected_colcount))
}
```

```
# Import "rawdata_347.txt" for "net migration rate"

# Read the data file
file_path <- paste0(working_directory_path, "/data/rawdata_347.txt")
lines <- readLines(file_path)

# Convert lines to a tibble
migration_rate <- map_dfr(lines, function(line) {
  parts <- strsplit(trimws(line), "\\s{2,}")[[1]]
  tibble(
```

```r
    Country = parts[2],
    Net_Migration_Rate = as.numeric(parts[3])
  )
}, .id = NULL) %>%
  filter(!is.na(Country), Country != "")

# Make sure all rows have been read
check_number_of_rows(227, 2, migration_rate)
head(migration_rate)
```

```
## # A tibble: 6 x 2
##   Country               Net_Migration_Rate
##   <chr>                              <dbl>
## 1 Syria                               27.1
## 2 British Virgin Islands              15.5
## 3 Luxembourg                          13.3
## 4 Cayman Islands                      13
## 5 Singapore                           11.8
## 6 Anguilla                            11.1
```

```r
# Import "rawdata_343.txt" for "median age"

# Read the data file
file_path <- paste0(working_directory_path, "/data/rawdata_343.txt")
lines <- readLines(file_path)

# Convert lines to a tibble
median_age <- map_dfr(lines, function(line) {
  parts <- strsplit(trimws(line), "\\s{2,}")[[1]]
  tibble(
    Country = parts[2],
    Median_Age = as.numeric(parts[3])
  )
}, .id = NULL) %>%
  filter(!is.na(Country), Country != "")

# Make sure all rows have been read
check_number_of_rows(227, 2, median_age)
head(median_age)
```

```
## # A tibble: 6 x 2
##   Country                   Median_Age
##   <chr>                          <dbl>
## 1 Monaco                          55.4
## 2 Japan                           48.6
## 3 Saint Pierre and Miquelon       48.5
## 4 Germany                         47.8
## 5 Italy                           46.5
## 6 Andorra                         46.2
```

```r
# Importing rawdata_373.csv (youth unemployment rate per country)
file_path <- paste0(working_directory_path, "/data/rawdata_373.csv")
```

```
youth_unemployment <- read_csv(file_path,
                               skip = 1, # Skip the predefined column names
                               col_names = c("Country", "Youth_Unemployment_Rate"),
                               col_types = c("c", "d"))  %>%
                    filter(!is.na(Country), Country != "")

# Make sure all rows have been read
check_number_of_rows(181, 2, youth_unemployment)
head(youth_unemployment)
```

```
## # A tibble: 6 x 2
##   Country           Youth_Unemployment_Rate
##   <chr>                               <dbl>
## 1 French Polynesia                     56.7
## 2 Kosovo                               55.4
## 3 South Africa                         53.4
## 4 Libya                                48.7
## 5 Eswatini                             47.1
## 6 Saint Lucia                          46.2
```

b.

Merge the data sets containing raw data using dplyr function on the unique keys. Keep the union of all
observations in the tables. What key are you using for merging? Return the dimension of the merged data
set.

Answer

```
# Merge the tibbles on the "Country" key
merged_country_data <- migration_rate %>%
  full_join(median_age, by = "Country") %>%
  full_join(youth_unemployment, by = "Country")


print(paste0("Dimensions: "))
```

```
## [1] "Dimensions: "
```

```
print(dim(merged_country_data))
```

```
## [1] 227   4
```

```
# Make sure the merge is correct
check_number_of_rows(227, 4, merged_country_data)
head(merged_country_data)
```

```
## # A tibble: 6 x 4
##   Country                Net_Migration_Rate Median_Age Youth_Unemployment_Rate
##   <chr>                               <dbl>      <dbl>                   <dbl>
## 1 Syria                                27.1       23.5                    35.8
## 2 British Virgin Islands               15.5       37.2                      NA
```

```
## 3 Luxembourg                   13.3      39.5                 14.2
## 4 Cayman Islands               13        40.5                 13.8
## 5 Singapore                    11.8      35.6                  9.1
## 6 Anguilla                     11.1      35.7                   NA
```

```r
# empty values
na_value_countries <- merged_country_data %>%
  filter(apply(., 1, anyNA))

print(nrow(na_value_countries))
```

```
## [1] 46
```

```r
head(na_value_countries)
```

```
## # A tibble: 6 x 4
##   Country            Net_Migration_Rate Median_Age Youth_Unemployment_Rate
##   <chr>                           <dbl>      <dbl>                   <dbl>
## 1 British Virgin Islands           15.5       37.2                      NA
## 2 Anguilla                         11.1       35.7                      NA
## 3 Turks and Caicos Islands          8.9       34.6                      NA
## 4 Aruba                             8.4       39.9                      NA
## 5 Sint Maarten                      6         41.1                      NA
## 6 Djibouti                          5.1       24.9                      NA
```

As expected, there are 46 rows with NA as value in the youth unemployment rate column, since the given
rawdata file, has fewer countries listed

c.

You will acquire more country level information such as the classification of the country based on income.
Such an information can be found at https://datahelpdesk.worldbank.org/knowledgebase/articles/906519.
From there extract the classification for 2020 into low/lower-middle/upper-middle/high income countries.

Answer

```r
# Importing rawdata from historical data file
file_path <- paste0(working_directory_path, "/data/OGHIST.xlsx")
full_excel_file <- read_excel(file_path,
                sheet = "Country Analytical History",
                range = cell_cols("A:AL"))
```

```r
head(full_excel_file)
```

```
## # A tibble: 6 x 38
##   ...1  World Bank Analytical ~1 ...3  ...4  ...5  ...6  ...7  ...8  ...9  ...10
##   <chr> <chr>                    <chr> <chr> <chr> <chr> <chr> <chr> <chr> <chr>
## 1 <NA>  (presented in World Dev~ <NA>  <NA>  <NA>  <NA>  <NA>  <NA>  <NA>  <NA>
## 2 <NA>  GNI per capita in US$ (~ <NA>  <NA>  <NA>  <NA>  <NA>  <NA>  <NA>  <NA>
## 3 <NA>  <NA>                     <NA>  <NA>  <NA>  <NA>  <NA>  <NA>  <NA>  <NA>
## 4 <NA>  Bank's fiscal year:      FY89  FY90  FY91  FY92  FY93  FY94  FY95  FY96
```

```
## 5 <NA>  Data for calendar year : 1987  1988   1989   1990   1991   1992   1993   1994
## 6 <NA>  Low income (L)            <= 4~ <= 5~ <= 5~ <= 6~ <= 6~ <= 6~ <= 6~ <= 7~
## # i abbreviated name: 1: 'World Bank Analytical Classifications'
## # i 28 more variables: ...11 <chr>, ...12 <chr>, ...13 <chr>, ...14 <chr>,
## #   ...15 <chr>, ...16 <chr>, ...17 <chr>, ...18 <chr>, ...19 <chr>,
## #   ...20 <chr>, ...21 <chr>, ...22 <chr>, ...23 <chr>, ...24 <chr>,
## #   ...25 <chr>, ...26 <chr>, ...27 <chr>, ...28 <chr>, ...29 <chr>,
## #   ...30 <chr>, ...31 <chr>, ...32 <chr>, ...33 <chr>, ...34 <chr>,
## #   ...35 <chr>, ...36 <chr>, ...37 <chr>, ...38 <chr>
```

```r
# Get the range for the year columns
year_col_range <- 3:ncol(full_excel_file)

# Extract the years and change column names
classification <- full_excel_file[11:nrow(full_excel_file), ]
colnames(classification) <-
  c("ISO", "Country", full_excel_file[5, year_col_range])

# Convert classification columns to factors, replacing ".." with NA
classification[, year_col_range] <-
  lapply(classification[, year_col_range],
         function(x) as.factor(replace(x, x == "..", NA)))

# Filter out rows where ISO is NA
classification <- classification[!is.na(classification$ISO), ]

head(classification)
```

```
## # A tibble: 6 x 38
##   ISO   Country   '1987' '1988' '1989' '1990' '1991' '1992' '1993' '1994' '1995'
##   <chr> <chr>     <fct>  <fct>  <fct>  <fct>  <fct>  <fct>  <fct>  <fct>  <fct>
## 1 AFG   Afghanis~ L      L      L      L      L      L      L      L      L
## 2 ALB   Albania   <NA>   <NA>   <NA>   LM     LM     LM     L      L      L
## 3 DZA   Algeria   UM     UM     LM     LM     LM     LM     LM     LM     LM
## 4 ASM   American~ H      H      H      UM     UM     UM     UM     UM     UM
## 5 AND   Andorra   <NA>   <NA>   <NA>   H      H      H      H      H      H
## 6 AGO   Angola    <NA>   LM     LM     LM     LM     LM     LM     LM     L
## # i 27 more variables: '1996' <fct>, '1997' <fct>, '1998' <fct>, '1999' <fct>,
## #   '2000' <fct>, '2001' <fct>, '2002' <fct>, '2003' <fct>, '2004' <fct>,
## #   '2005' <fct>, '2006' <fct>, '2007' <fct>, '2008' <fct>, '2009' <fct>,
## #   '2010' <fct>, '2011' <fct>, '2012' <fct>, '2013' <fct>, '2014' <fct>,
## #   '2015' <fct>, '2016' <fct>, '2017' <fct>, '2018' <fct>, '2019' <fct>,
## #   '2020' <fct>, '2021' <fct>, '2022' <fct>
```

```r
# Get classification for 2020
classification_2020 <- classification[, c("ISO", "Country", "2020")]
colnames(classification_2020) <- c("ISO", "Country", "Classification_2020")
print(classification_2020[order(classification_2020$Country), ])
```

```
## # A tibble: 224 x 3
##    ISO   Country         Classification_2020
##    <chr> <chr>           <fct>
## 1 AFG   Afghanistan     L
```

```
##  2 ALB    Albania               UM
##  3 DZA    Algeria               LM
##  4 ASM    American Samoa        UM
##  5 AND    Andorra               H
##  6 AGO    Angola                LM
##  7 ATG    Antigua and Barbuda H
##  8 ARG    Argentina             UM
##  9 ARM    Armenia               UM
## 10 ABW    Aruba                 H
## # i 214 more rows
```

d.

Merge this information to the data set in b.

1. What are the common variables? Can you merge using them? Why or why not?

2. A reliable merging for countries are ISO codes as they are standardized across data sources. Download the mapping of ISO codes to countries from https://www.cia.gov/the-world-factbook/references/countrydata-codes/ and load it

3. Merge the data sets using the ISO codes.

Answer

```
# Check for countries which are in my merged list,
# but not in the classification list,
# if just merged by country name
missing_countries <- setdiff(merged_country_data$Country, classification_2020$Country)
print(sort(missing_countries))
```

```
##  [1] "Anguilla"
##  [2] "Brunei"
##  [3] "Burma"
##  [4] "Congo, Democratic Republic of the"
##  [5] "Congo, Republic of the"
##  [6] "Cook Islands"
##  [7] "Cote d'Ivoire"
##  [8] "Curacao"
##  [9] "Czechia"
## [10] "Egypt"
## [11] "Faroe Islands"
## [12] "Gaza Strip"
## [13] "Guernsey"
## [14] "Hong Kong"
## [15] "Iran"
## [16] "Jersey"
## [17] "Korea, North"
## [18] "Korea, South"
## [19] "Kyrgyzstan"
## [20] "Laos"
## [21] "Macau"
## [22] "Macedonia"
```

```
## [23] "Micronesia, Federated States of"
## [24] "Montserrat"
## [25] "Russia"
## [26] "Saint Barthelemy"
## [27] "Saint Helena, Ascension, and Tristan da Cunha"
## [28] "Saint Kitts and Nevis"
## [29] "Saint Lucia"
## [30] "Saint Martin"
## [31] "Saint Pierre and Miquelon"
## [32] "Saint Vincent and the Grenadines"
## [33] "Sao Tome and Principe"
## [34] "Sint Maarten"
## [35] "Slovakia"
## [36] "Syria"
## [37] "Taiwan"
## [38] "Turkey"
## [39] "Venezuela"
## [40] "Virgin Islands"
## [41] "Wallis and Futuna"
## [42] "West Bank"
## [43] "Yemen"
```

If we would just merge by the country name, there would be over 40 countries missing, which are in my merged_data list. The reason could be e.g. different spelling, different order (Korea, South) or the countries are just not included. In summary, a lack of standardization hinders us in linking the data

```r
# Importing country data codes from "Country Data Codes"
file_path <- paste0(working_directory_path, "/data/Country Data Codes.csv")
country_data_codes <- read_csv(file_path, show_col_types = FALSE)

# Get subset and rename columns
iso <- country_data_codes[, c("Name", "GENC")]
colnames(iso) <- c("Country", "ISO")
iso[iso == "-"] <- NA

# Merge iso into existing data set
merged_country_data_with_iso <- merged_country_data %>%
  full_join(iso, by = "Country")
merged_country_data <- merged_country_data_with_iso

head(merged_country_data)
```

```
## # A tibble: 6 x 5
##   Country         Net_Migration_Rate Median_Age Youth_Unemployment_R~1 ISO
##   <chr>                        <dbl>      <dbl>                  <dbl> <chr>
## 1 Syria                         27.1       23.5                   35.8 SYR
## 2 British Virgin Isl~           15.5       37.2                   NA   VGB
## 3 Luxembourg                    13.3       39.5                   14.2 LUX
## 4 Cayman Islands                13         40.5                   13.8 CYM
## 5 Singapore                     11.8       35.6                    9.1 SGP
## 6 Anguilla                      11.1       35.7                   NA   AIA
## # i abbreviated name: 1: Youth_Unemployment_Rate
```

Even though we added most codes, there are still a few missing, which have to be added manually since the matching is not perfect

```
# List countries without iso
print(merged_country_data[is.na(merged_country_data$ISO), ])
```

```
## # A tibble: 10 x 5
##    Country          Net_Migration_Rate Median_Age Youth_Unemployment_R~1 ISO
##    <chr>                         <dbl>      <dbl>                  <dbl> <chr>
##  1 Macedonia                       0.4         39                   45.4 <NA>
##  2 Turkey                         -4.3       32.2                   20.2 <NA>
##  3 France, Metropoli~               NA         NA                     NA <NA>
##  4 Myanmar                          NA         NA                     NA <NA>
##  5 United States Min~               NA         NA                     NA <NA>
##  6 Virgin Islands (U~               NA         NA                     NA <NA>
##  7 Virgin Islands (U~               NA         NA                     NA <NA>
##  8 Western Samoa                    NA         NA                     NA <NA>
##  9 World                            NA         NA                     NA <NA>
## 10 Zaire                            NA         NA                     NA <NA>
## # i abbreviated name: 1: Youth_Unemployment_Rate
```

```
merged_country_data$ISO[merged_country_data$Country == "Turkey"] <- "TUR"
merged_country_data$ISO[merged_country_data$Country == "Macedonia"] <- "MKD"

# List countries without ISO
print(merged_country_data[is.na(merged_country_data$ISO), ])
```

```
## # A tibble: 8 x 5
##   Country          Net_Migration_Rate Median_Age Youth_Unemployment_R~1 ISO
##   <chr>                         <dbl>      <dbl>                  <dbl> <chr>
## 1 France, Metropolit~              NA         NA                     NA <NA>
## 2 Myanmar                          NA         NA                     NA <NA>
## 3 United States Mino~              NA         NA                     NA <NA>
## 4 Virgin Islands (UK)              NA         NA                     NA <NA>
## 5 Virgin Islands (US)              NA         NA                     NA <NA>
## 6 Western Samoa                    NA         NA                     NA <NA>
## 7 World                            NA         NA                     NA <NA>
## 8 Zaire                            NA         NA                     NA <NA>
## # i abbreviated name: 1: Youth_Unemployment_Rate
```

For the countries that are still without ISO, there are special political and regional reasons, which is why they cannot be added.

```
# Merge classification into existing data set
merged_country_data_with_class_2020 <- merged_country_data %>%
  full_join(classification_2020[, c("ISO", "Classification_2020")], by = "ISO")
merged_country_data <- merged_country_data_with_class_2020

head(merged_country_data)
```

```
## # A tibble: 6 x 6
##   Country          Net_Migration_Rate Median_Age Youth_Unemployment_R~1 ISO
```

```
##   <chr>                             <dbl>   <dbl>            <dbl> <chr>
## 1 Syria                              27.1    23.5            35.8 SYR
## 2 British Virgin Isl~                15.5    37.2            NA    VGB
## 3 Luxembourg                         13.3    39.5            14.2 LUX
## 4 Cayman Islands                     13      40.5            13.8 CYM
## 5 Singapore                          11.8    35.6             9.1 SGP
## 6 Anguilla                           11.1    35.7            NA    AIA
## # i abbreviated name: 1: Youth_Unemployment_Rate
## # i 1 more variable: Classification_2020 <fct>
```

e.

Introduce into the data set information on continent for each country and subcontinent (region). You should find a way to gather this data. You can find an appropriate online resource, download the data and merge the information with the existing data set. Name the merged data set df_vars.

Answer

To add the requested region data, the following dataset has been used: https://statisticstimes.com/geography/countries-by-continents.php

```r
# Importing continent and region data
file_path <- paste0(working_directory_path, "/data/continent_region_data.csv")
continent_region_data <- read_delim(file_path,
                                    delim = ";",
                                    locale = locale(encoding = "UTF-8"),
                                    show_col_types = FALSE)

# Get subset and rename columns
continent_region_data_subset <-
  continent_region_data[, c("ISO-alpha3", "Region 1", "Continent")]
colnames(continent_region_data_subset) <-
 c("ISO", "Region", "Continent")

# Merge into existing data set
merged_country_data_with_region <- merged_country_data %>%
  full_join(continent_region_data_subset, by = "ISO")

# Create new dataset
df_vars <- merged_country_data_with_region
head(df_vars)
```

```
## # A tibble: 6 x 8
##    Country          Net_Migration_Rate Median_Age Youth_Unemployment_R~1 ISO
##    <chr>                         <dbl>      <dbl>                  <dbl> <chr>
## 1 Syria                          27.1       23.5                   35.8 SYR
## 2 British Virgin Isl~            15.5       37.2                   NA    VGB
## 3 Luxembourg                     13.3       39.5                   14.2 LUX
## 4 Cayman Islands                 13         40.5                   13.8 CYM
## 5 Singapore                      11.8       35.6                    9.1 SGP
## 6 Anguilla                       11.1       35.7                   NA    AIA
## # i abbreviated name: 1: Youth_Unemployment_Rate
## # i 3 more variables: Classification_2020 <fct>, Region <chr>, Continent <chr>
```

f.

Discuss on the tidyness of the data set df_vars. What are the observational units, what are the variables?
What can be considered fixed vs measured variables? Tidy the data if needed.

Answer

```r
str(df_vars)
```

```
## tibble [293 x 8] (S3: tbl_df/tbl/data.frame)
##  $ Country               : chr [1:293] "Syria" "British Virgin Islands" "Luxembourg" "Cayman Islands
##  $ Net_Migration_Rate    : num [1:293] 27.1 15.5 13.3 13 11.8 11.1 10.6 8.9 8.4 8.3 ...
##  $ Median_Age            : num [1:293] 23.5 37.2 39.5 40.5 35.6 35.7 32.9 34.6 39.9 55.4 ...
##  $ Youth_Unemployment_Rate: num [1:293] 35.8 NA 14.2 13.8 9.1 NA 5.3 NA NA 26.6 ...
##  $ ISO                   : chr [1:293] "SYR" "VGB" "LUX" "CYM" ...
##  $ Classification_2020   : Factor w/ 4 levels "H","L","LM","UM": 2 1 1 1 1 NA 1 1 1 1 ...
##  $ Region                : chr [1:293] "Western Asia" "Caribbean" "Western Europe" "Caribbean" ...
##  $ Continent             : chr [1:293] "Asia" "North America" "Europe" "North America" ...
```

```r
# save the data for further usage
write_csv(df_vars, path = "../data/df_vars.csv")
```

```
## Warning: The 'path' argument of 'write_csv()' is deprecated as of readr 1.4.0.
## i Please use the 'file' argument instead.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_lifecycle_warnings()' to see where this warning was
## generated.
```

Observational units vs variables*

In the given data frame (df_vars), the observational units are the countries. Each row represents a different
country.

The variables in the data frame are:

1. `Country`: The name of the country (character/string variable).
2. `Net_Migration_Rate`: The net migration rate for the country (numeric variable).
3. `Median_Age`: The median age of the population in the country (numeric variable).
4. `Youth_Unemployment_Rate`: The youth unemployment rate in the country (numeric variable).
5. `ISO`: The ISO code for the country (character/string variable).
6. `Classification_2020`: The classification of the country, which is a factor with four levels: "H", "L",
   "LM", and "UM" (factor variable).
7. `Region`: The region to which the country belongs (character/string variable).
8. `Continent`: The continent where the country is located (character/string variable).

as can be seen via str(df_vars)

Fixed vs mesured variables

Fixed variables are those that do not vary within the dataset. In this case, the fixed variables are: - `Country`
- `ISO` - `Region` - `Continent`

These variables are essentially identifiers or labels for the observational units (countries) and do not change
within the dataset.

Measured variables are those that are measured for each observational unit. In this case, the measured variables are: - `Net_Migration_Rate` - `Median_Age` - `Youth_Unemployment_Rate` - `Classification_2020`

These variables represent quantitative or categorical measurements or characteristics of the countries.

Tydiness

A tidy dataset has the following characteristics:

1. Each variable forms a column.
2. Each observation forms a row.
3. Each type of observational unit forms a table.

Analyzing df_vars:

1. Each variable forms a column:

   - Each variable (`Country`, `Net_Migration_Rate`, ...) forms a separate column, which satisfies this principle.

2. Each observation forms a row:

   - In this case, each row represents an observation, which is a country. Therefore, this principle is also satisfied.

3. Each type of observational unit forms a table:

   - The dataset contains only one type of observational unit, which is countries. All the observations (rows) and variables (columns) pertain to countries, so this principle is satisfied as well.

Based on this analysis, the dataset appears to be in a tidy format.

2. Data analysis - Part 1

g.

Make a frequency table for the status variable in the merged data set. Briefly comment on the results.

```
df_vars %>% mutate(status = forcats::fct_infreq(Classification_2020)) %>% count(status)
```

```
##    status  n
## 1       H 80
## 2      UM 57
## 3      LM 55
## 4       L 27
## 5    <NA> 74
```

We can see that the biggest part of the countries actually falls into the high income class, in the middle classes "UM" and "LM" there are about an even amount of countries. Only 27 countries fall under low income. However there are relatively many NA values present which might skew the previous observations.

h.

What is the distribution of income status in the different continents? Compute the absolute frequencies as well as the relative frequency of status within each continent. Briefly comment on the results.

```
by_cont <- df_vars %>% group_by(Continent)
income_status <- by_cont %>% count(Classification_2020) %>% mutate(abs = n, freq = n/sum(n))
income_status$n <- NULL
income_status
```

```
## # A tibble: 29 x 4
## # Groups:   Continent [8]
##    Continent  Classification_2020   abs   freq
##    <chr>      <chr>               <int>  <dbl>
##  1 Africa     H                       1 0.0167
##  2 Africa     L                      23 0.383
##  3 Africa     LM                     23 0.383
##  4 Africa     UM                      7 0.117
##  5 Africa     <NA>                    6 0.1
##  6 Antarctica <NA>                    1 1
##  7 Asia       H                      14 0.275
##  8 Asia       L                       4 0.0784
##  9 Asia       LM                     19 0.373
## 10 Asia       UM                     14 0.275
## # i 19 more rows
```

From this table it is clearly visible that the continents with the highes income are Europe and North America, both of which have no countries classified in "L" and very few in "LM". Next is Oceania also having no "L"-countries, but an even spread between the other categories. South America is also void of "L"-classified countries, with most of them being in the "UM" category. Asia has a few "L"-countries and also an even spread between the others. The poorest continent is Africa, only having one country in the "H" category and many in the "L" and "LM" categories.

i.

From h. identify the countries which are the only ones in their respective group. Explain in few words the output.

```
abs_1 <- income_status %>% filter(abs==1) %>% select(Continent, Classification_2020)
abs_1 <- left_join(abs_1, df_vars)
```

```
## Joining with `by = join_by(Continent, Classification_2020)`
```

```
abs_1 %>% select(Continent, Classification_2020, Country)
```

```
## # A tibble: 5 x 3
## # Groups:   Continent [5]
##   Continent     Classification_2020 Country
##   <chr>         <chr>               <chr>
## 1 Africa        H                   Seychelles
## 2 Antarctica    <NA>                Antarctica
## 3 Europe        LM                  Ukraine
## 4 South America LM                  Bolivia
## 5 <NA>          UM                  <NA>
```

In the output we have 1 NA value, Antarctica (which is the only country on the continent), Seychelles (the only "H" rated country in Africa) and Ukraine and Bolivia which are the only "LM" rated countries in Europe/South America respectively.

j.

For each continent count the number of sub-regions in the data set. How granular are the subcontinents that you employ in the analysis?

```
by_cont %>% summarise(sub_regions = n_distinct(Region))
```

```
## # A tibble: 8 x 2
##   Continent      sub_regions
##   <chr>                <int>
## 1 Africa                   5
## 2 Antarctica               1
## 3 Asia                     5
## 4 Europe                   4
## 5 North America            3
## 6 Oceania                  4
## 7 South America            1
## 8 <NA>                     1
```

The continents are separated into sub regions in a very different way. None of them have more than 5 sub-regions, however while a rather small continent such as Europe has 4 sub-regions, really big ones like Asia and Africa are also only split into 5. South America even has only 1 sub-region.

k.

Look at the frequency distribution of income status in the subregions of North- and South-Americas. Comment on the results.

```
na <- df_vars %>% filter(Continent=="North America")
sa <- df_vars %>% filter(Continent=="South America")
americas <- full_join(na, sa)
```

```
## Joining with `by = join_by(Country, Net_Migration_Rate, Median_Age,
## Youth_Unemployment_Rate, ISO, Classification_2020, Region, Continent)`
```

```
by_region <- americas %>% group_by(Region)
income_status_am <- by_region %>% count(Classification_2020) %>% mutate(abs = n, freq = n/sum(n))
income_status_am$n <- NULL
income_status_am
```

```
## # A tibble: 12 x 4
## # Groups:   Region [4]
##   Region          Classification_2020   abs   freq
##   <chr>           <chr>               <int>  <dbl>
## 1 Caribbean       H                      14 0.5
## 2 Caribbean       LM                      1 0.0357
## 3 Caribbean       UM                      7 0.25
## 4 Caribbean       <NA>                    6 0.214
## 5 Central America LM                      4 0.5
## 6 Central America UM                      4 0.5
## 7 Northern America H                      4 0.8
```

```
##  8 Northern America <NA>                     1 0.2
##  9 South America    H                         2 0.125
## 10 South America    LM                        1 0.0625
## 11 South America    UM                        8 0.5
## 12 South America    <NA>                      5 0.312
```

The Caribbean seems to have rather high income with the biggest part of countries being classified in "H", the same is true for North America. Central America is equally split between "LM" and "UM". South American countries are mostly classified in "UM", with a few in "H" and "LM".

l.

Dig deeper into the low-middle income countries of the Americas. Which ones are they? Are they primarily small island states in the Caribbean? Comment.

```
americas_LM <- income_status_am %>% filter(Classification_2020=="LM")
americas_LM <- left_join(americas_LM, df_vars)
```

```
## Joining with 'by = join_by(Region, Classification_2020)'
```

```
americas_LM
```

```
## # A tibble: 6 x 10
## # Groups:   Region [3]
##   Region  Classification_2020   abs   freq Country Net_Migration_Rate Median_Age
##   <chr>   <chr>               <int>  <dbl> <chr>                <dbl>      <dbl>
## 1 Caribb~ LM                      1 0.0357 Haiti                 -1.9       24.1
## 2 Centra~ LM                      4 0.5    Belize                 -1       23.9
## 3 Centra~ LM                      4 0.5    Hondur~               -1.4       24.4
## 4 Centra~ LM                      4 0.5    Nicara~               -2.4       27.3
## 5 Centra~ LM                      4 0.5    El Sal~               -4.8       27.7
## 6 South ~ LM                      1 0.0625 Bolivia               -0.3       25.3
## # i 3 more variables: Youth_Unemployment_Rate <dbl>, ISO <chr>, Continent <chr>
```

They are primarily states in Central America, only one of them is a Central American island nation.

3. Data analysis - Part 2

```
df_vars <- read.csv("../data/df_vars.csv")
df_vars$Classification_2020 <- factor(
  df_vars$Classification_2020,
  levels = c("L", "LM", "UM", "H")
  )
```

m.

Create a table of average values for median age, youth unemployment rate and net migration rate separated into income status. Make sure that in the output, the ordering of the income classes is proper (i.e., L, LM, UM, H or the other way around). Briefly comment the results.

Answer

```r
average_values <- df_vars %>%
  group_by(Classification_2020) %>%
  arrange(Classification_2020) %>%
  summarize(
    Avg_Net_Migration_rate = mean(Net_Migration_Rate, na.rm = TRUE),
    Avg_Median_Age = mean(Median_Age, na.rm = TRUE),
    Avg_Youth_Unemployment_rate = mean(Youth_Unemployment_Rate, na.rm = TRUE)
  )
print(average_values)
```

```
## # A tibble: 5 x 4
##   Classification_2020 Avg_Net_Migration_rate Avg_Median_Age
##   <fct>                                 <dbl>          <dbl>
## 1 L                                    -0.519           19.2
## 2 LM                                   -1.92            24.8
## 3 UM                                   -4.32            32.0
## 4 H                                     1.89            39.4
## 5 <NA>                                 -3.36            35.6
## # i 1 more variable: Avg_Youth_Unemployment_rate <dbl>
```

Regarding the average net migration rate, it is interesting to note that all income classes feature negative values except for the higher income class. A negative migration rate suggests that more people are emigrating than immigrating in the country, which could point towards some form of "Brain drain", which of course would have some effect on the countries economy.\ The average median age is also highest in higher income countries, which reminds of often cited statistics linking economic prosperity with a lower birth rate and generally older populations.

n.

Look also at the standard deviation instead of the mean in m. Do you gain additional insights? Briefly comment the results.

Answer

```r
average_standard_deviation <- df_vars %>%
  group_by(Classification_2020) %>%
  arrange(Classification_2020) %>%
  summarize(
    Avg_Net_Migration_rate = sd(Net_Migration_Rate, na.rm = TRUE),
    Avg_Median_Age = sd(Median_Age, na.rm = TRUE),
    Avg_Youth_Unemployment_rate = sd(Youth_Unemployment_Rate, na.rm = TRUE)
  )
print(average_standard_deviation)
```

```
## # A tibble: 5 x 4
##   Classification_2020 Avg_Net_Migration_rate Avg_Median_Age
##   <fct>                                 <dbl>          <dbl>
## 1 L                                      6.03           3.64
## 2 LM                                     3.49           4.93
## 3 UM                                    12.8            6.35
## 4 H                                      5.69           5.43
## 5 <NA>                                   9.22           8.99
## # i 1 more variable: Avg_Youth_Unemployment_rate <dbl>
```

Standard deviation is a measure spread in the underlying distribution of data. Regarding sampled data, it can be seen as a measure of "insecurity" in the data. Thus, the standard deviation of the average net migration rate of the upper middle income class immediately stands out. In a more thorough data analysis, I would suggest performing a deeper investigation of the values in this class. Also, outlier detection could be performed or more robust measures than standard deviation, such as MAD or Trimmed Standard Deviation could be used.

o.

Repeat the analysis in m. for each income status and continent combination. Discuss the results.

Answer

```
average_values_continent <- df_vars %>%
  group_by(Classification_2020, Continent) %>%
  arrange(Classification_2020) %>%
  summarize(
    Avg_Net_Migration_rate = mean(Net_Migration_Rate, na.rm = TRUE),
    Avg_Median_Age = mean(Median_Age, na.rm = TRUE),
    Avg_Youth_Unemployment_rate = mean(Youth_Unemployment_Rate, na.rm = TRUE)
  )
```

```
## 'summarise()' has grouped output by 'Classification_2020'. You can override
## using the '.groups' argument.
```

```
print(average_values_continent)
```

```
## # A tibble: 29 x 5
## # Groups:   Classification_2020 [5]
##    Classification_2020 Continent     Avg_Net_Migration_rate Avg_Median_Age
##    <fct>               <chr>                          <dbl>          <dbl>
##  1 L                   Africa                         -1.77           18.3
##  2 L                   Asia                            6.7            24.4
##  3 LM                  Africa                         -1.41           21.8
##  4 LM                  Asia                           -1.51           27.6
##  5 LM                  Europe                          2.3            41.2
##  6 LM                  North America                  -2.3            25.5
##  7 LM                  Oceania                        -5.78           24.7
##  8 LM                  South America                  -0.3            25.3
##  9 UM                  Africa                          0.843          25.6
## 10 UM                  Asia                           -9.61           31.9
## # i 19 more rows
## # i 1 more variable: Avg_Youth_Unemployment_rate <dbl>
```

Naturally, a grouping of every income class and continent pairing returns a lot more rows than the queries before. While it is hard to point out much given this list, it is intriguing for example to look at the average median age for the contintent Europe. It is common to think of Europe as a continent with an old population, and this also shows here, with Europe having the oldest average median age for all income classes compared with other continents.

p.

Identify countries which are doing well in terms of both youth unemployment and net migration rate (in the top 25% of their respective continent in terms of net migration rate and in the bottom 25% of their respective continent in terms of youth unemployment).

Answer

```r
continent_quartiles <- df_vars %>%
  group_by(Continent) %>%
  summarize(
    Net_Migration_Q1 = quantile(Net_Migration_Rate, 0.25, na.rm = TRUE),
    Youth_Unemployment_Q3 = quantile(Youth_Unemployment_Rate, 0.75, na.rm = TRUE)
  )

countries_in_top25_net_migration <- df_vars %>%
  inner_join(continent_quartiles, by = "Continent") %>%
  filter(
    Net_Migration_Rate >= Net_Migration_Q1,
    Youth_Unemployment_Rate <= Youth_Unemployment_Q3
  ) %>%
  select(Country)

print(countries_in_top25_net_migration)
```

```
##                        Country
## 1                   Luxembourg
## 2                Cayman Islands
## 3                    Singapore
## 4                      Bahrain
## 5                    Australia
## 6                  New Zealand
## 7                       Cyprus
## 8         United Arab Emirates
## 9                        Malta
## 10                       Qatar
## 11                      Canada
## 12                  Isle of Man
## 13                      Sweden
## 14                     Belgium
## 15                 Switzerland
## 16                      Norway
## 17                     Ireland
## 18                     Austria
## 19                     Iceland
## 20                       Macau
## 21               United States
## 22                     Denmark
## 23                     Finland
## 24              United Kingdom
## 25                     Czechia
## 26                 Korea, South
## 27                     Ukraine
## 28                      Israel
```

```
## 29                              Netherlands
## 30                               Hong Kong
## 31                                  Russia
## 32                                 Germany
## 33                                Slovenia
## 34                                 Hungary
## 35                           Cote d'Ivoire
## 36                                  France
## 37                              Seychelles
## 38                                   Palau
## 39                              Costa Rica
## 40                                 Belarus
## 41                                 Suriname
## 42                              Kazakhstan
## 43                                   Benin
## 44                                   Chile
## 45                                Portugal
## 46                                Slovakia
## 47                              Azerbaijan
## 48                             Bahamas, The
## 49                                  Bhutan
## 50                                 Ecuador
## 51                                  Guinea
## 52                                  Guyana
## 53                                   Japan
## 54                              Madagascar
## 55                               Mauritius
## 56                        Papua New Guinea
## 57                                Thailand
## 58                                    Togo
## 59                                  Zambia
## 60                             Afghanistan
## 61                                  Panama
## 62                                Paraguay
## 63                                Ethiopia
## 64                                   Kenya
## 65                                 Nigeria
## 66                                 Bolivia
## 67                                Cambodia
## 68                                Cameroon
## 69                                   Egypt
## 70                                Malaysia
## 71                                 Vietnam
## 72                                    Oman
## 73                                Tanzania
## 74                            Burkina Faso
## 75                              Cabo Verde
## 76                                Colombia
## 77                                   Niger
## 78                                 Burundi
## 79                              Mauritania
## 80                                Mongolia
## 81  Congo, Democratic Republic of the
## 82                                Pakistan
```

```
## 83                     Belize
## 84                       Laos
## 85                  Indonesia
## 86               Sierra Leone
## 87                    Senegal
## 88                  Sri Lanka
## 89                      Burma
## 90                   Honduras
## 91               Gambia, The
## 92                      Ghana
## 93            Solomon Islands
## 94                  Guatemala
## 95                     Mexico
## 96                  Nicaragua
## 97         Dominican Republic
## 98                   Kiribati
## 99                       Cuba
## 100                 West Bank
## 101          Marshall Islands
## 102                      Fiji
## 103                    Tuvalu
## 104                     Nauru
```

r.

Export the final data set to a csv with ";" separator and "." as a symbol for missing values; no rownames should be included in the csv. Upload the .csv to TUWEL together with your .Rmd and .html (or .pdf).

Answer

```r
write.table(
  df_vars,
  file = "../data/final_data.csv",
  sep = ";",
  na = ".",
  row.names = FALSE
  )
```